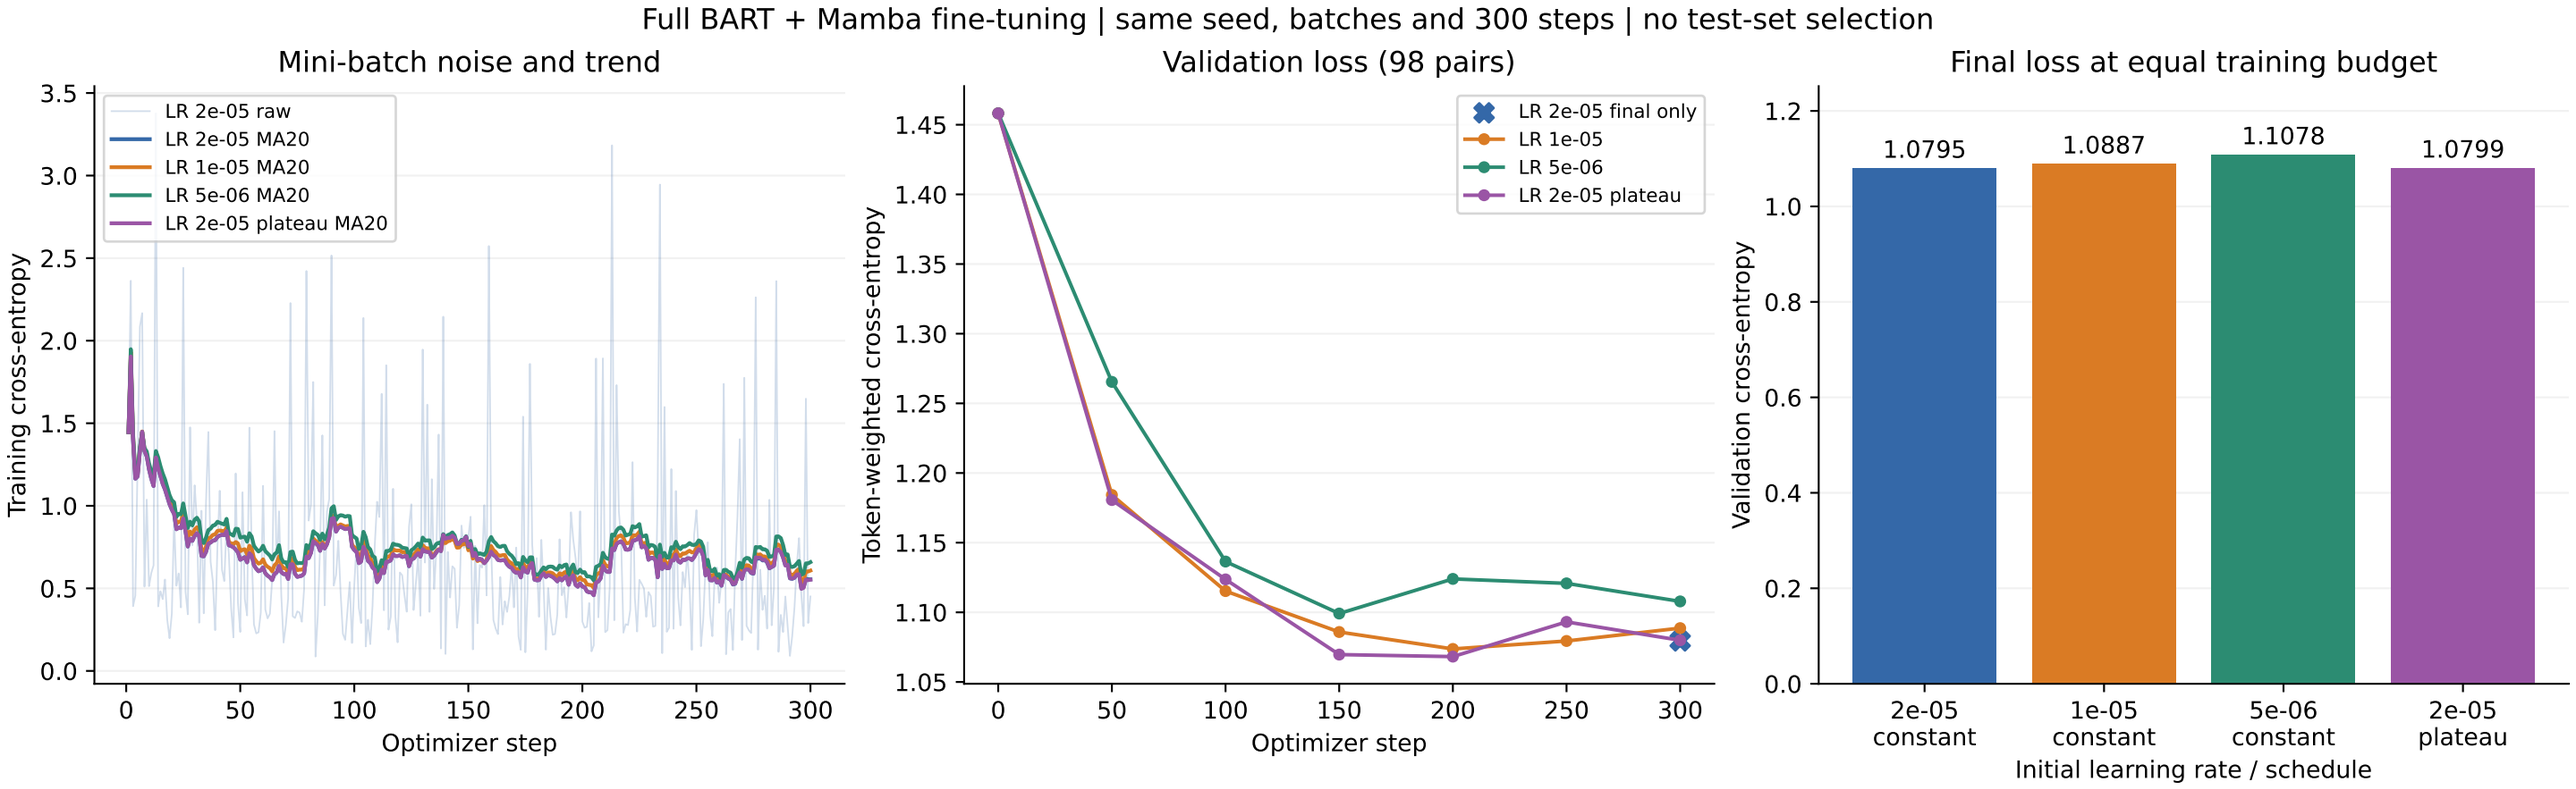
Full BART + Mamba fine-tuning | same seed, batches and 300 steps | no test-set selection
Mini-batch noise and trend
Training cross-entropy
Optimizer step
LR 2e-05 raw
LR 2e-05 MA20
LR 1e-05 MA20
LR 5e-06 MA20
LR 2e-05 plateau MA20
Validation loss (98 pairs)
Token-weighted cross-entropy
Optimizer step
LR 2e-05 final only
LR 1e-05
LR 5e-06
LR 2e-05 plateau
Final loss at equal training budget
Validation cross-entropy
1.0795
1.0887
1.1078
1.0799
2e-05 constant
1e-05 constant
5e-06 constant
2e-05 plateau
Initial learning rate / schedule